


# Introduction to Operating Systems

CPSC/ECE 3220 Summer 2018

Lecture Notes  
OSPP Chapter 13

(adapted by Mark Smotherman from Tom Anderson's slides on OSPP web site)

# Named Data in a File System



- Directory – typically tree-structured
- Index structure – typically tree-structured
- Free space map – often a bitmap
- Locality heuristics
  - Policy in finding free space (first-fit, etc.)
  - Grouping of directories and files
  - Defragmentation
  - Optimization of writes over reads

# Main Points

- File layout
- Directory layout
- Access control

# File Organization Design Constraints

- For small files:
  - Small blocks for storage efficiency
  - Files used together should be stored together
- For large files:
  - Contiguous allocation for sequential access
  - Efficient lookup for random access
- May not know at file creation
  - Whether file will become small or large

# File System Design

- Data structures
  - Directories: file name -> file metadata
    - Store directories as files
  - File metadata: how to find file data blocks
  - Free map: list of free disk blocks
- How do we organize these data structures?
  - Device has non-uniform performance

# Design Challenges

- Index structure
  - How do we locate the blocks of a file?
- Index granularity
  - What block size do we use?
- Free space
  - How do we find unused blocks on disk?
- Locality
  - How do we preserve spatial locality?
- Reliability
  - What if machine crashes in middle of a file system op?

# File Organization

- Want sequential data placement that provides efficient sequential access
- Also want placement that provides efficient random access
- But...
  - Reject contiguous storage of disk blocks
    - why?
  - Reject linked-list storage with links located in the disk blocks – why?

# File Organization (2)

- Typically tree-structured indexing of data
  - Want to limit overheads to be efficient for small files
  - Provide scalability for large files
  - Provide a place for per-file metadata



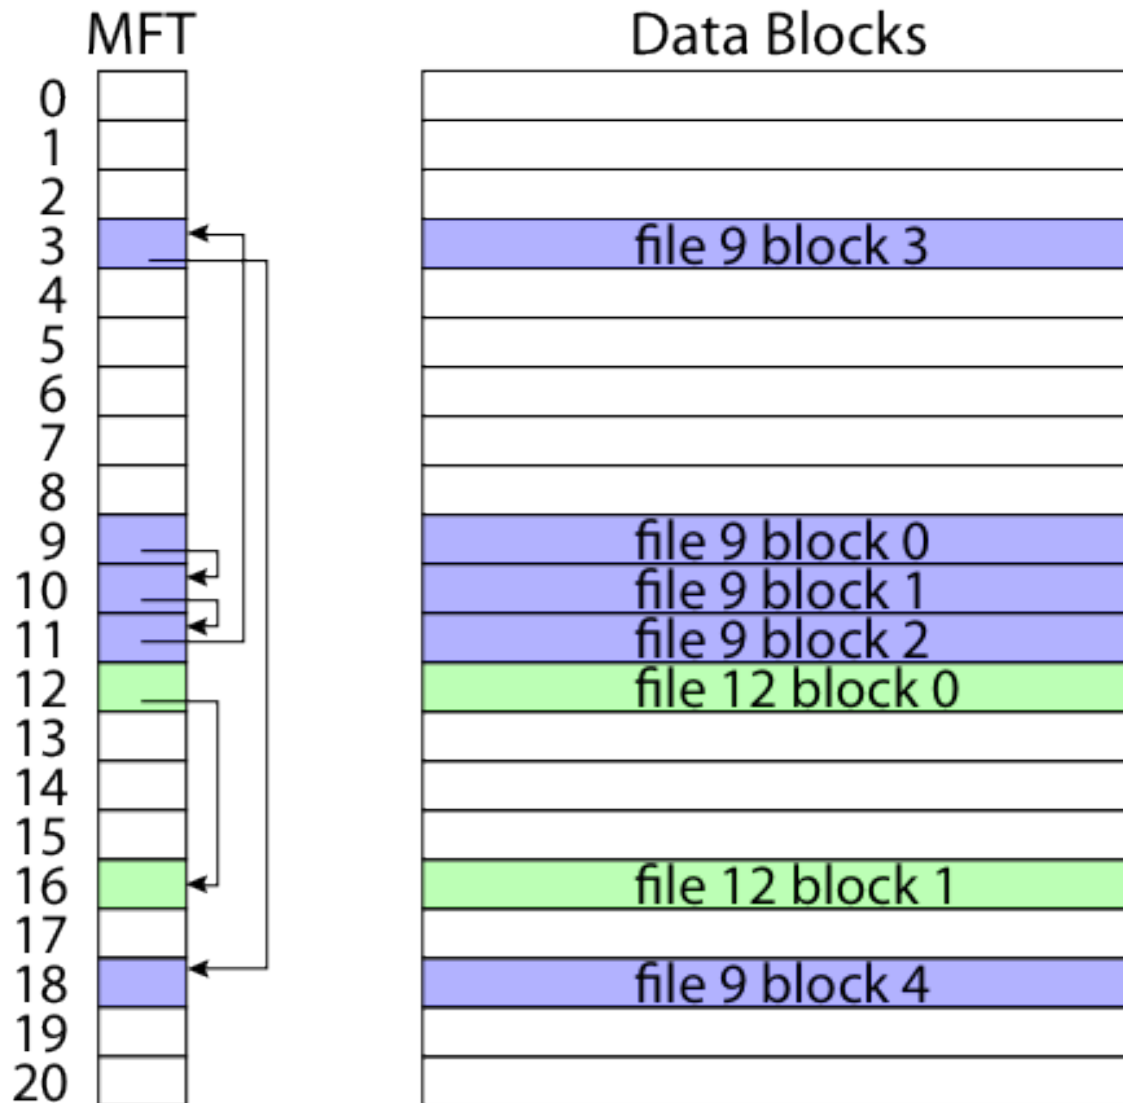
# File Organization Design Options

	<b>FAT</b>	<b>FFS</b>	<b>NTFS</b>	<b>ZFS</b>
<b>Index structure</b>	linked list	tree (fixed, asymmetric)	tree (dynamic)	tree (COW, dynamic)
<b>Granularity</b>	block	block	extent	block
<b>Free space allocation</b>	FAT array	bitmap (fixed location)	bitmap (file)	space map (log-structured)
<b>Locality</b>	defrag	block groups, first fit, reserve space	extents, best fit, defrag	write-anywhere block groups

# Microsoft File Allocation Table (FAT)

- Linked list index structure
  - Links are collected together in a table, not placed out in data blocks
  - Simple, easy to implement
  - Still widely used (e.g., thumb drives)
- File table
  - Linear map of all blocks on disk
  - Each file a linked list of blocks

# FAT Organization



# FAT Evaluation

- Pros:
  - Easy to find free block
  - Easy to append to a file
  - Easy to delete a file
- Cons:
  - Small file access is slow
  - Random access is very slow
  - Fragmentation
    - File blocks for a given file may be scattered
    - Files in the same directory may be scattered
    - Problem becomes worse as disk fills

# Berkeley UNIX FFS (Fast File System)

- inode table
  - Analogous to FAT table
- inode
  - Metadata
    - File owner, access permissions, access times, ...
  - Set of 12 data pointers
  - With 4KB blocks => max size of 48KB files

# Additional FFS inode Block Pointers

- Indirect block pointer (13<sup>th</sup> pointer in inode)
  - pointer to disk block of data pointers
  - 4KB block size => 1K data blocks => 4MB
- Doubly indirect block pointer (14<sup>th</sup> in inode)
  - Doubly indirect block => 1K indirect blocks
  - 4GB (+ 4MB + 48KB)
- Triply indirect block pointer (15<sup>th</sup> in inode)
  - Triply indirect block => 1K doubly indirect blocks
  - 4TB (+ 4GB + 4MB + 48KB)

# FFS inode

File Metadata
Direct Pointer
DP
DP
DP
DP
DP
DP
DP
DP
DP
DP
Direct Pointer
Indirect Pointer
Dbl. Indirect Ptr.
Tripl. Indirect Ptr.


# FFS Asymmetric Tree

- Small files: shallow tree
  - Efficient storage for small files
- Large files: deep tree
  - Efficient lookup for random access in large files
- Sparse files: only fill pointers if needed




# FFS Small File


Inode Array



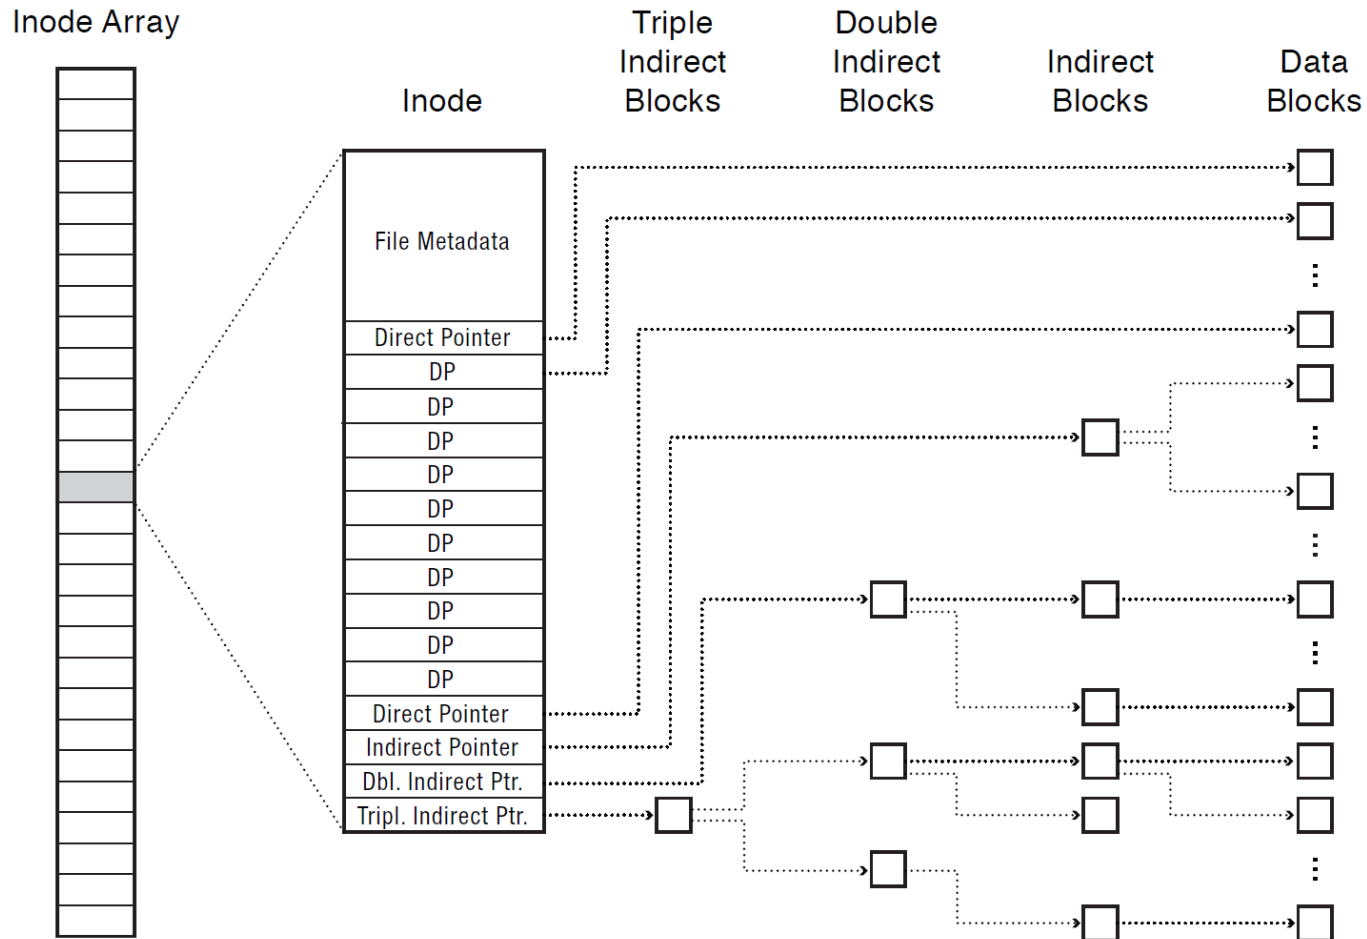
Inode



Data Blocks

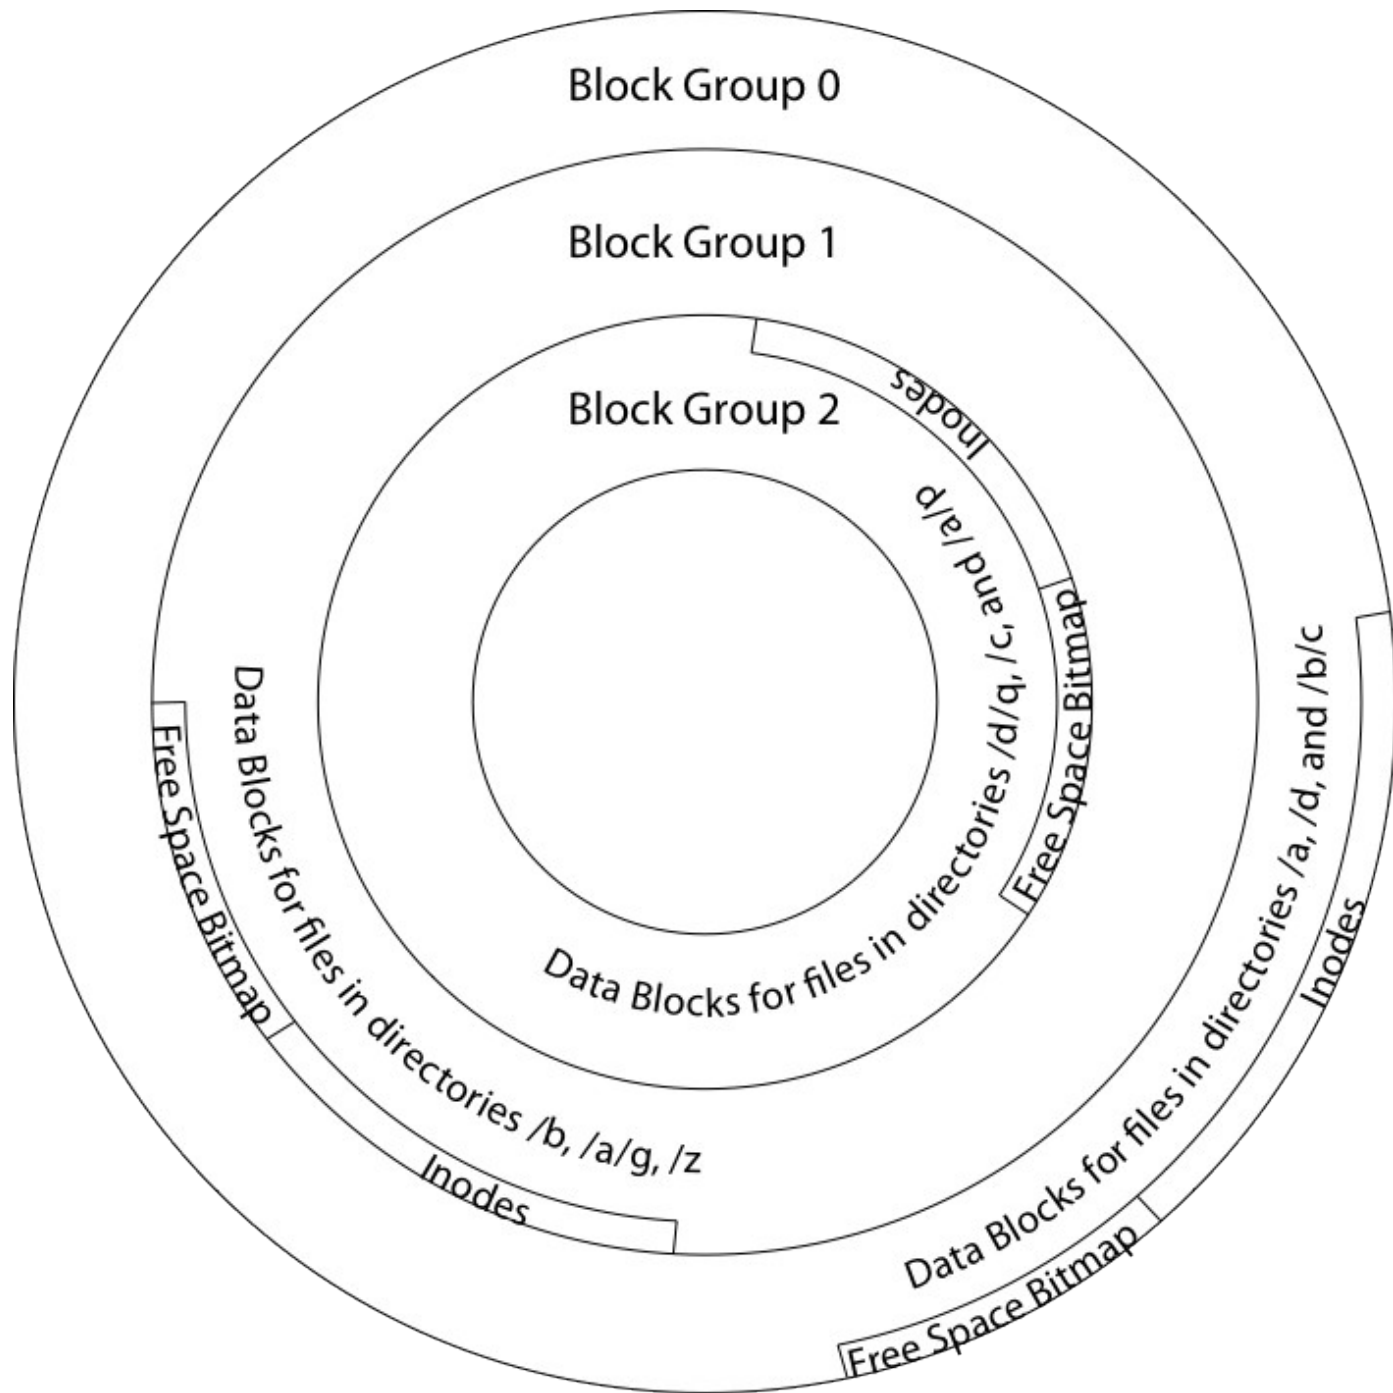


# FFS Large File




# FFS Locality


- Block group allocation
  - Block group is a set of nearby cylinders
  - Files in same directory located in same group
  - Subdirectories located in different block groups
- inode table spread throughout disk
  - inodes, bitmap near file blocks
- First fit allocation
  - Small files fragmented, large files contiguous




# FFS First Fit Block Allocation



# FFS First Fit Block Allocation



# FFS First Fit Block Allocation



# FFS Evaluation

- Pros
  - Efficient storage for both small and large files
  - Locality for both small and large files
  - Locality for metadata and data
- Cons
  - Inefficient for tiny files (a 1 byte file requires both an inode and a data block)
  - Inefficient encoding when file is mostly contiguous on disk (no equivalent to superpages)
  - Need to reserve 10-20% of free space to prevent fragmentation

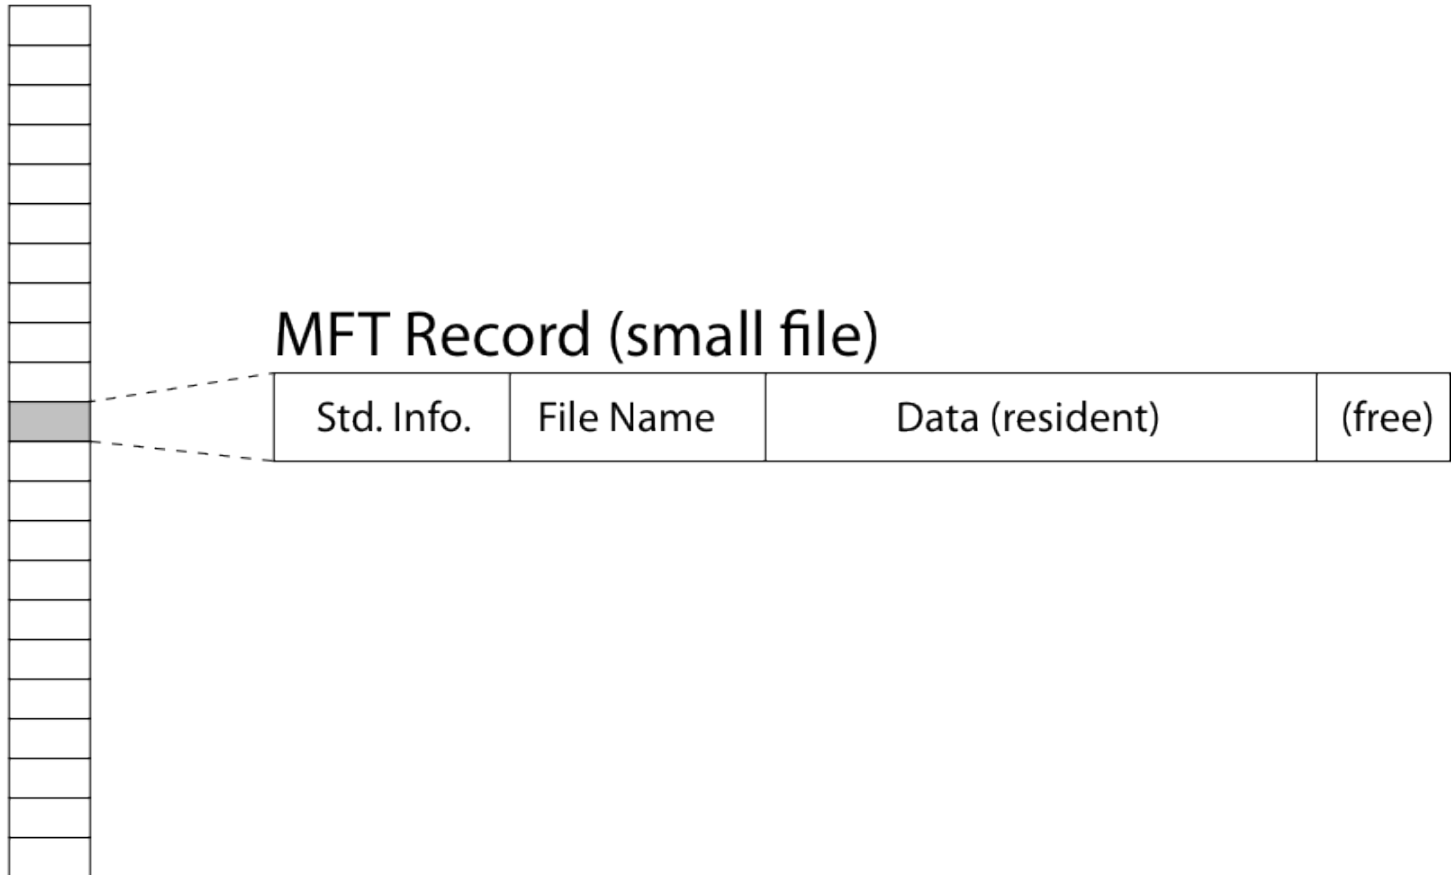


# NTFS

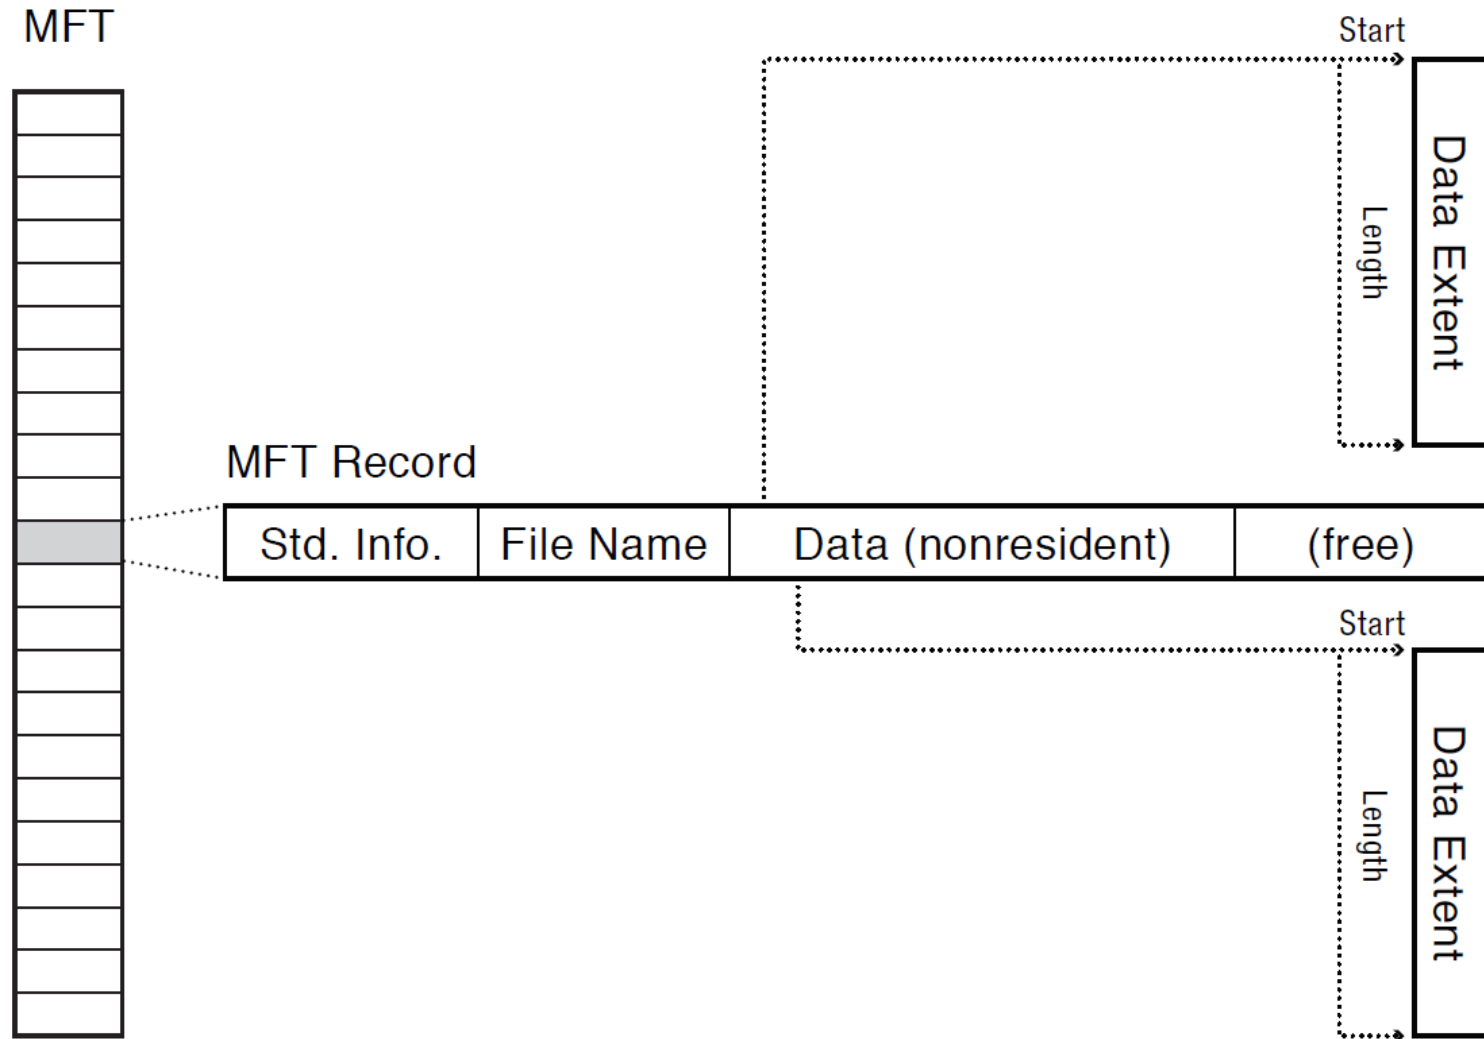
- Master File Table
  - Flexible 1KB storage for metadata and data
- Extents
  - Block pointers cover runs of blocks
  - Similar approach in linux (ext4)
  - File create can provide hint as to size of file
- Journalling for reliability
  - Next chapter

# NTFS Small File

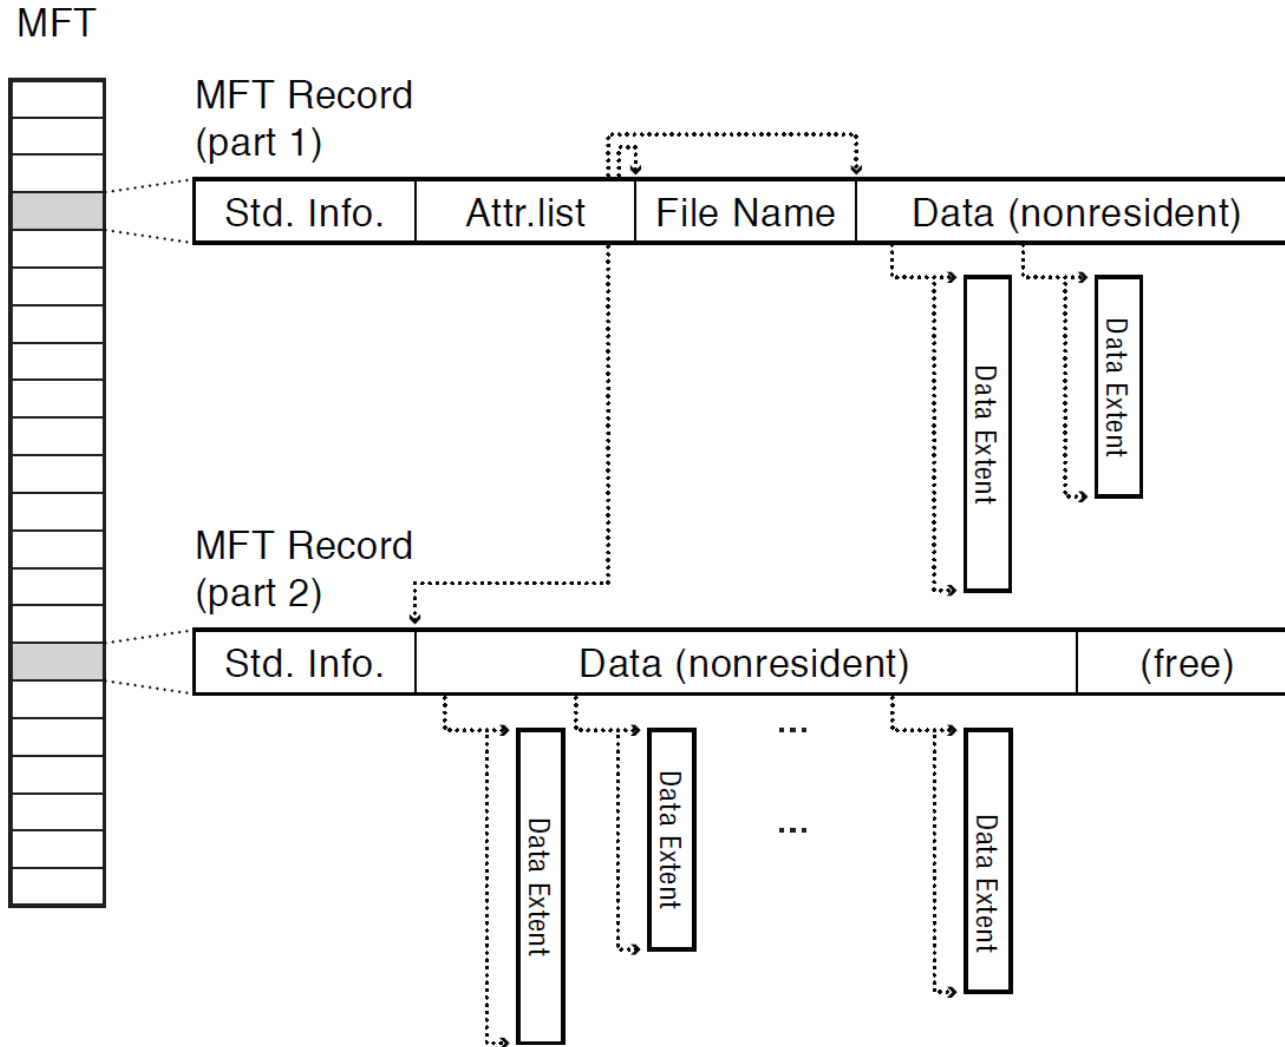
# Master File Table



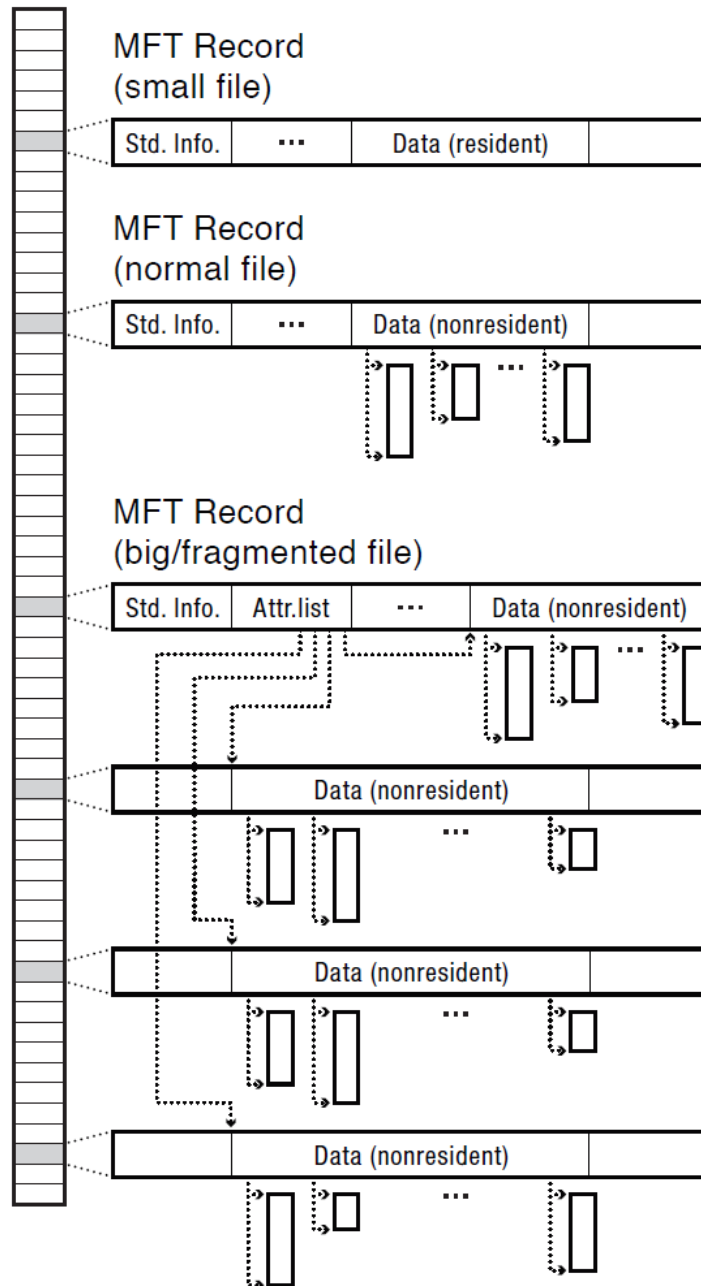
# NTFS Medium-Sized File



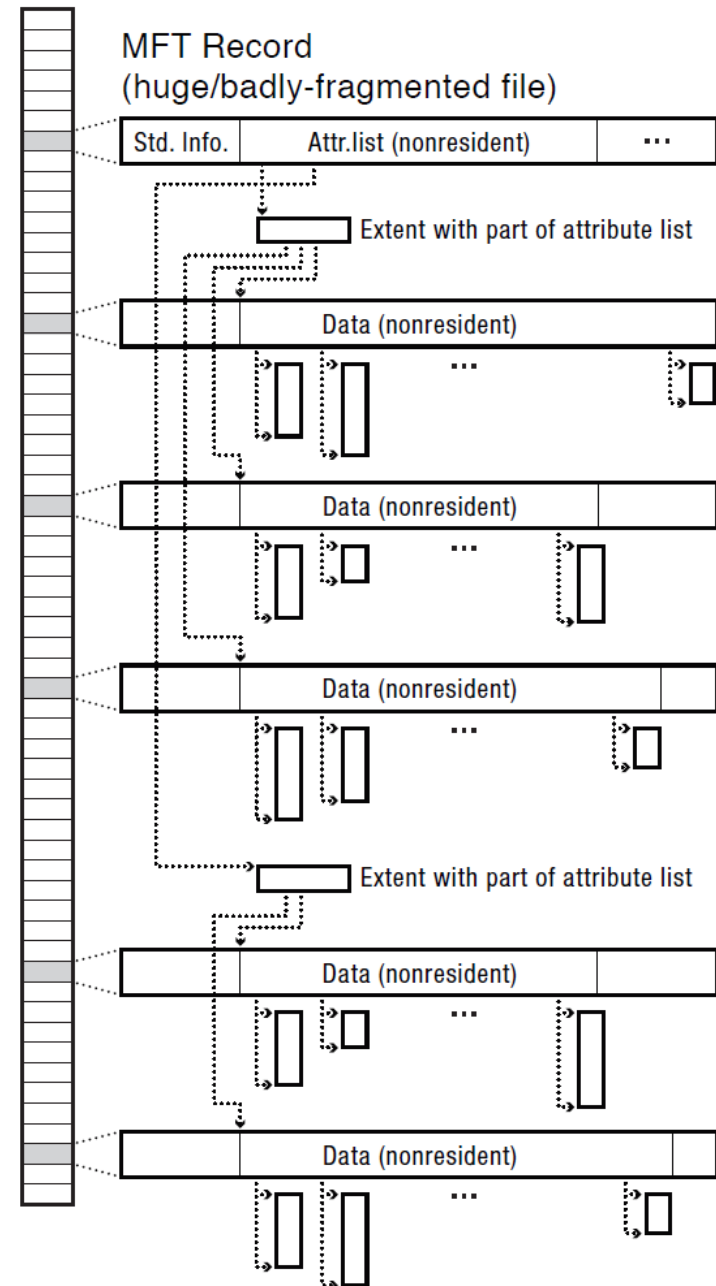
# NTFS Indirect Block




MFT



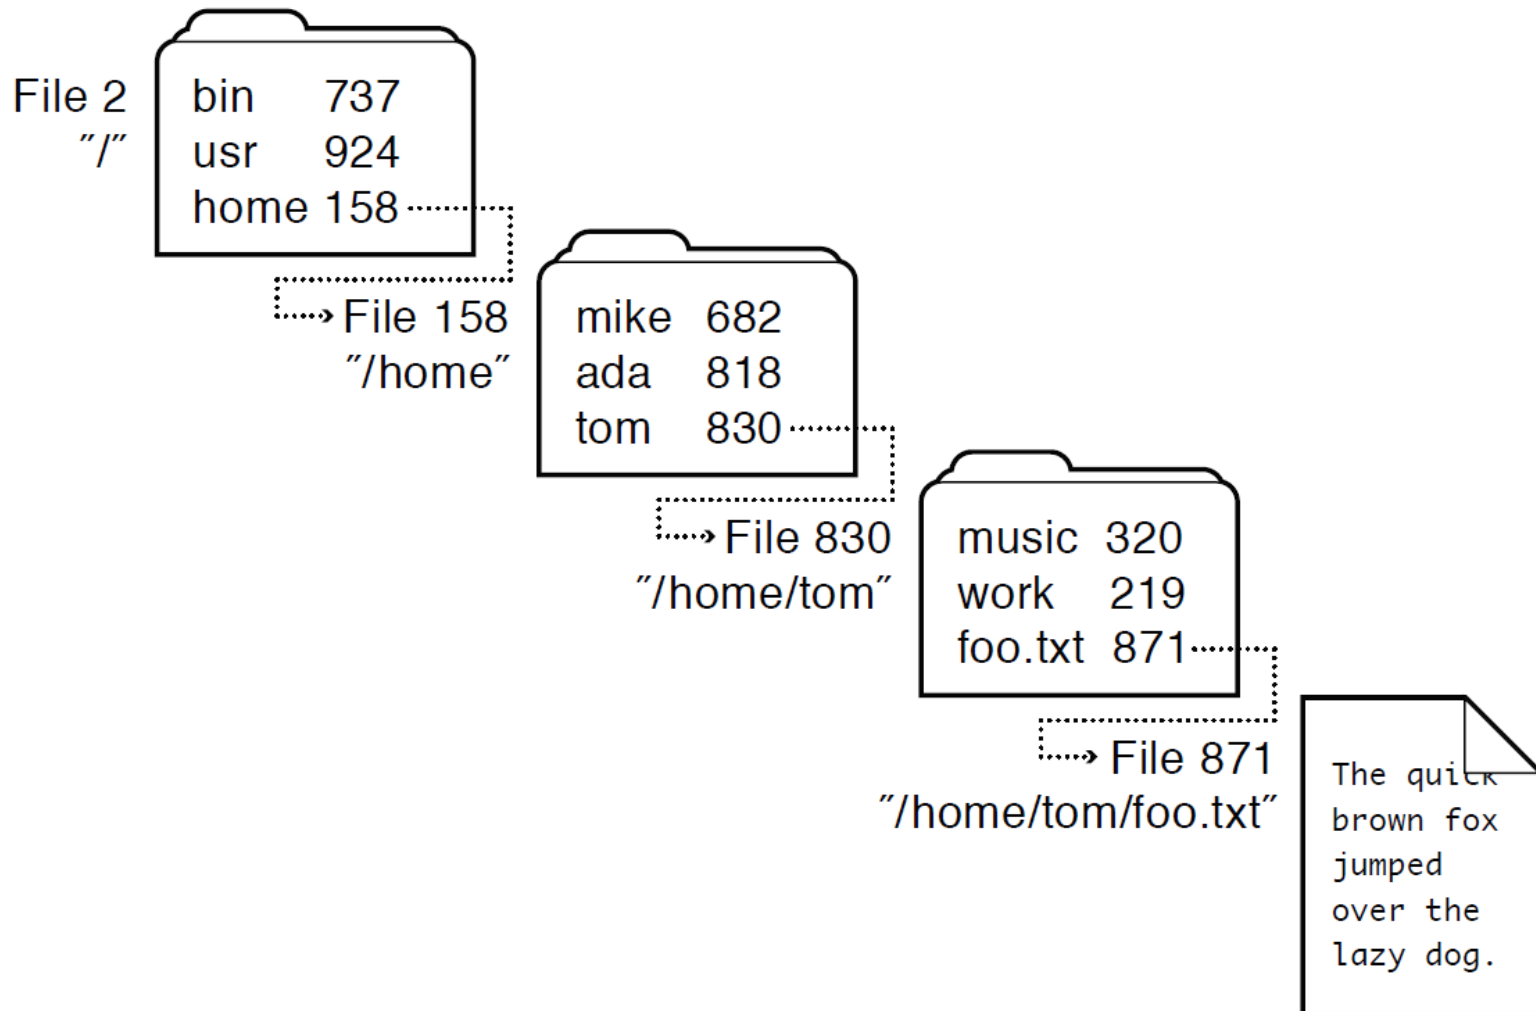
MFT



# Directories Are Files



# Recursive Filename Lookup



# Directory Layout

Directory stored as a file  
Linear search to find filename (small directories)

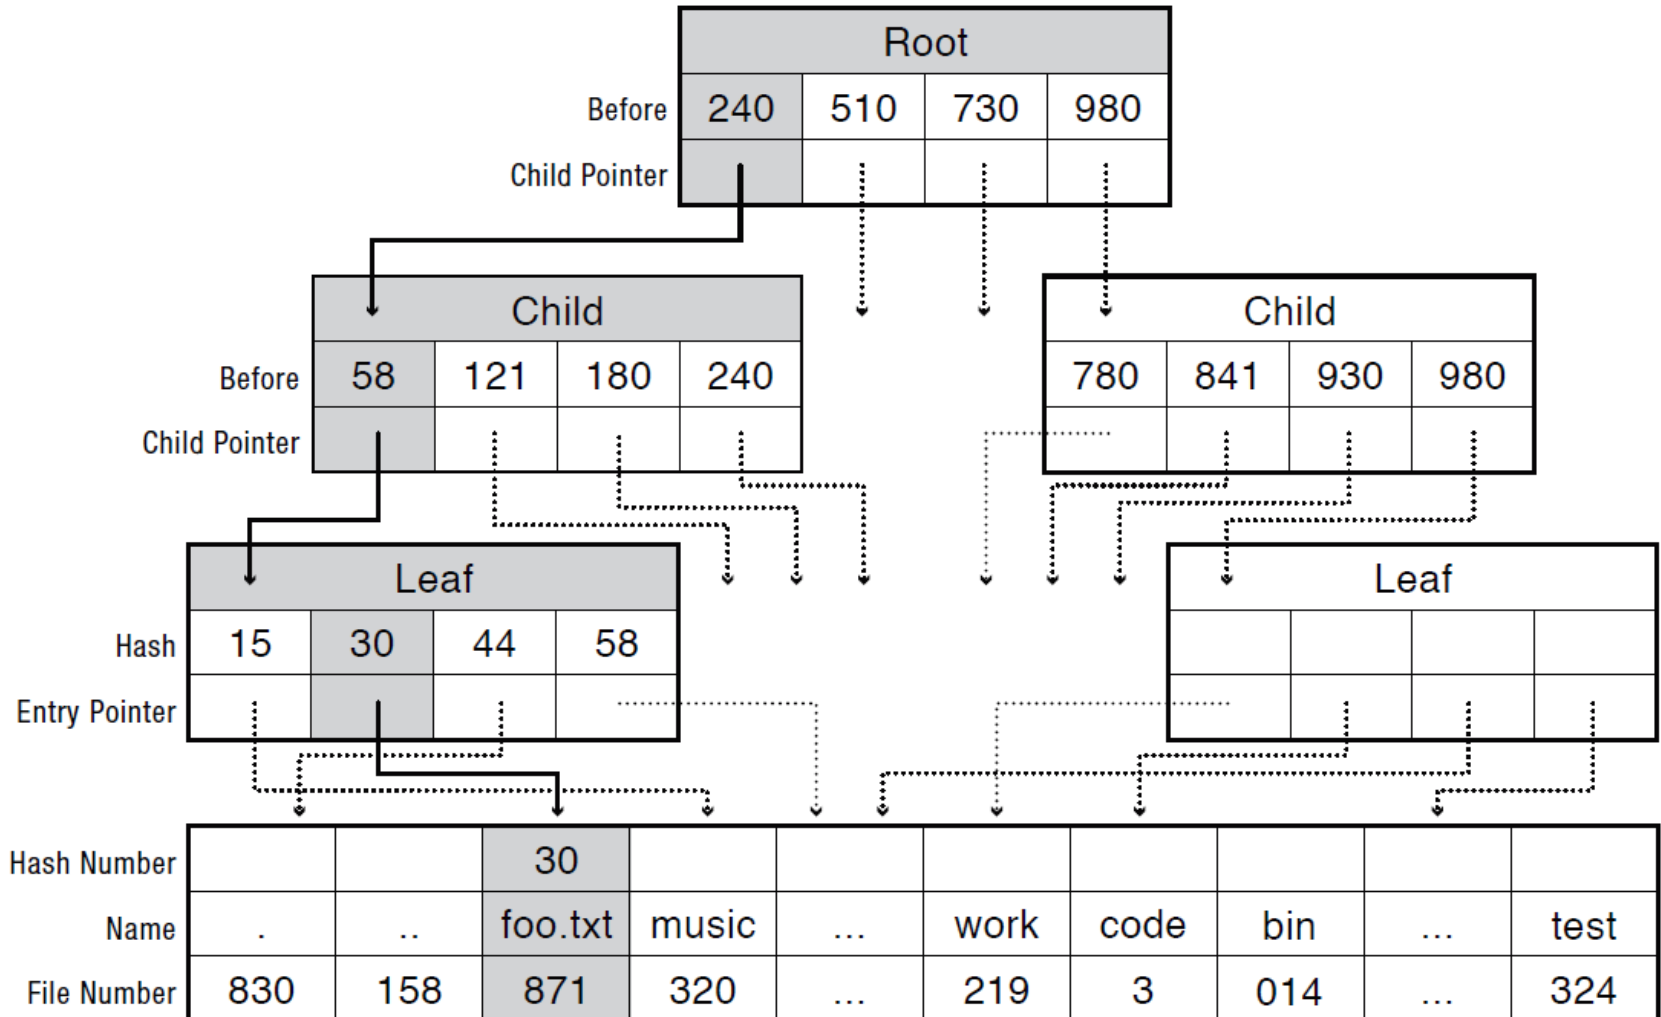
File 830  
"/home/tom"

Name	.	..	music	work	Free Space	foo.txt	Free Space	End of File
File Number	830	158	320	219		871		
Next	↓	↓	↓	↓		↓		
	↓	↓	↓	↓	↓	↓	↓	



# Large Directories: B Trees

Search for Hash (foo.txt) = 0x30



# Large Directories: Layout

## File Containing Directory

Name	...	music	work	...	...	Root	Child	Leaf	Leaf	Child	...
File Number		320	219								
	Directory Entries					B+Tree Nodes					

# Access Control

- Access control matrix
  - Row for each user (or application)
  - Column for every protected resource
  - Sparse
- Row-oriented approach
  - Capability list
  - Keep with user
- Column-oriented approach
  - Access control list (ACL)
  - Keep with resource

# Capability List

- Example: process keeps a table of open files with allowed type of access
- Example: smartphone app keeps a list of resources with allowed type of access

# Access Control List

- For every protected resource, list of who is permitted to do what
- Example: Windows ACL
  - List of access control entries (ACEs)
  - Each ACE contains a 32-bit access rights mask and a security identifier (SID)
- Example: compressed ACL for UNIX files
  - Owner, group, world: read, write, execute
  - Setuid: program will run using the permissions of user who installed it
  - File type to indicate if file is a directory