

# Introduction to Operating Systems

CPSC/ECE 3220 Summer 2018

Lecture Notes  
OSPP Chapter 5 – Part A

(adapted by Mark Smotherman from Tom Anderson's slides on OSPP web site)

# Lost Update / Record Out Problem

## Thread A with shared "x"

$x = x + 1;$

// at machine code level

load r1, x

add r1, r1, #1

// switch threads 

store r1, x

## Thread B with shared "x"

$x = x + 1;$

// at machine code level

load r1, x

add r1, r1, #1

store r1, x

// switch threads

# Question: Are All Pushes Successful?

Consider shared linked list accessed as a stack: (adapted from Michael Scott)

```
struct node{  
    ...  
    struct node *next;  
}  
  
struct node *head;  
  
void push( struct node *new ){  
    new->next = head;  
    // what if a second push starts at this point?  
    head = new;  
}
```

# Synchronization Motivation

- When threads concurrently read/write shared memory, program behavior is undefined
  - Two threads write to the same variable; which one should win?
- Thread schedule is non-deterministic
  - Behavior can change when you re-run program
- Also:
  - Compiler/hardware instruction reordering
  - Multi-word operations are not atomic

# Question: Can this panic?

## Thread 1

```
p =  
  someComputation(  
  );  
pInitalized = true;
```

## Thread 2

```
while ( !pInitalized )  
  /* empty loop  
  body */ ;  
q = someFunction(p);  
if ( q !=  
    someFunction(p) )  
  panic();
```

# Why Reordering?

- Why do compilers reorder instructions?
  - Efficient code generation requires analyzing control/data dependency
  - If variables can spontaneously change, most compiler optimizations become impossible
- Why do CPUs reorder instructions?
  - Write buffering: allow next instruction to execute while write is being completed

## Fix: **memory barrier**

- Instruction to compiler/CPU
- All ops before barrier complete before barrier returns
- No op after barrier starts until barrier returns

# Too Much Milk Example

	<u>Person A</u>	<u>Person B</u>
12:30	Look in fridge. Out of milk.	
12:35	Leave for store.	
12:40	Arrive at store.	Look in fridge. Out of milk.
12:45	Buy milk.	Leave for store.
12:50	Arrive home, put milk away.	Arrive at store.
12:55		Buy milk.
1:00		Arrive home, put milk away. Oh no!

# Definitions

**Race condition:** output of a concurrent program depends on the order of operations between threads

**Mutual exclusion:** only one thread does a particular thing at a time

- **Critical section:** piece of code that only one thread can execute at once

**Lock:** prevent someone from doing something

- Lock before entering critical section, before accessing shared data
- Unlock when leaving, after done accessing shared data
- Wait if locked (all synchronization involves waiting!)



# Too Much Milk, Try #1

- Correctness property
  - Someone buys if needed (liveness)
  - At most one person buys (safety)
- Try #1: leave a note

```
if ( !milk )
    if ( !note ) {
        leave note
        buy milk
        remove note
    }
```

# Too Much Milk, Try #2

## Thread A

```
leave note A
if ( !note B ) {
    if ( !milk )
        buy milk
}
remove note A
```

## Thread B

```
leave note B
if ( !noteA ) {
    if ( !milk )
        buy milk
}
remove note B
```

# Too Much Milk, Try #3

## Thread A

```
leave note A
while ( note B ) //
  X
  do nothing;
if ( !milk )
  buy milk;
remove note A
```

## Thread B

```
leave note B
if ( !noteA ) { // Y
  if ( !milk )
    buy milk
}
remove note B
```

Can guarantee at X and Y that  
either:

- (i) Safe for me to buy
- (ii) Other will buy, ok to quit

# Lessons

- Solution is complicated
  - “obvious” code often has bugs
- Modern compilers/architectures reorder instructions
  - Making reasoning even more difficult
- Generalizing to many threads/processors
  - Even more complex: see Peterson’s algorithm

# Roadmap

Concurrent Applications

---

Shared Objects

Bounded Buffer

Barrier

---

Synchronization Variables

Semaphores

Locks

Condition Variables

---

Atomic Instructions

Interrupt Disable

Test-and-Set

---

Hardware

Multiple Processors

Hardware Interrupts

---

# Locks

- `Lock::acquire()`
    - wait until lock is free, then take it
  - `Lock::release()`
    - release lock, waking up anyone waiting for it
1. At most one lock holder at a time (safety)
  2. If no one holding, acquire gets lock (progress)
  3. If all lock holders finish and no higher priority waiters, waiter eventually gets lock (progress)

# Question: Why only Acquire/Release

- Suppose we add a method to ask if a lock is free. Suppose it returns true. Is the lock:
  - Free?
  - Busy?
  - Don't know?

# Too Much Milk, #4

Locks allow concurrent code to be much simpler:

```
lock.acquire();
```

```
if (!milk)
```

```
    buy milk
```

```
lock.release();
```



# Lock Example: Malloc/Free

```
char *malloc (n) {  
    heaplock.acquire();  
    p = allocate  
    memory  
  
    heaplock.release();  
    return p;  
}
```

```
void free(char *p) {  
    heaplock.acquire();  
    put p back on  
    free list  
  
    heaplock.release();  
}
```

# Rules for Using Locks

- Lock is initially free
- Always acquire before accessing shared data structure
  - Beginning of procedure!
- Always release after finishing with shared data
  - End of procedure!
  - Only the lock holder can release
  - DO NOT throw lock for someone else to release
- Never access shared data without lock
  - Danger!

# Example: Bounded Buffer

```
tryget() {  
    item = NULL;  
    lock.acquire();  
    if ( front < tail ) {  
        item = buf[front %  
MAX];  
        front++; // ignoring ovf  
    }  
    lock.release();  
    return item;  
}
```

```
tryput(item) {  
    lock.acquire();  
    if ( (tail - front) <  
size ) {  
        buf[tail % MAX] =  
item;  
        tail++; // ignoring  
ovf  
    }  
    lock.release();  
}
```

ally: front = tail = 0; lock = FREE; MAX is buffer capacity

# Example: Bounded Buffer

- For simplicity, assume no wraparound on the integers front and tail; I'll assume you can fix that if you want
  - front = total number of items that have ever been removed
  - tail = total number of items ever inserted
- Lock at beginning of procedure; unlock at end; no access outside of locks
- Note that we don't know whether the buffer is still empty once we release the lock– we only know the state of the buffer while holding the lock!

# Question

- If `tryget()` returns `NULL`, do we know the buffer is empty?
- If we poll `tryget()` in a loop, what happens to a thread calling `tryput()`?

# Condition Variables

- Waiting inside a critical section
  - Called only when holding a lock
- Wait: atomically release lock and relinquish processor
  - Reacquire the lock when wakened
- Signal: wake up a waiter, if any
- Broadcast: wake up all waiters, if any

# Condition Variable Design Pattern

```
methodThatWaits() {  
    lock.acquire();  
    // Read/write shared state  
  
    while ( !  
testSharedState() ) {  
        cv.wait(&lock);  
    }  
  
    // Read/write shared state  
    lock.release();  
}
```

```
methodThatSignals() {  
    lock.acquire();  
    // Read/write shared state  
  
    // If testSharedState is now  
true  
    cv.signal(&lock);  
  
    // Read/write shared state  
    lock.release();  
}
```

# Example: Bounded Buffer

```
get() {  
    lock.acquire();  
    while ( front == tail ) {  
        empty.wait(&lock);  
    }  
    item = buf[front % MAX];  
    front++; // ignoring ovf  
    full.signal(&lock);  
    lock.release();  
    return item;  
}
```

```
put(item) {  
    lock.acquire();  
    while ( (tail - front) ==  
MAX ) {  
        full.wait(&lock);  
    }  
    buf[tail % MAX] = item;  
    tail++; // ignoring ovf  
    empty.signal(&lock);  
    lock.release();  
}
```

Initially: front = tail = 0; MAX is buffer capacity  
empty/full are condition variables



# Pre/Post Conditions

- What is state of the bounded buffer at lock acquire?
  - $\text{front} \leq \text{tail}$
  - $\text{front} + \text{MAX} \geq \text{tail}$
- These are also true on return from wait
- And at lock release
- Allows for proof of correctness

# Pre/Post Conditions

```
methodThatWaits() {  
    lock.acquire();  
    // Pre-condition: State is consistent  
  
    // Read/write shared state  
  
    while ( !testSharedState() ) {  
        cv.wait(&lock);  
    }  
    // WARNING: shared state may  
    // have changed! But  
    // testSharedState is TRUE  
    // and pre-condition is true  
  
    // Read/write shared state  
    lock.release();  
}
```

```
methodThatSignals() {  
    lock.acquire();  
    // Pre-condition: State is consistent  
  
    // Read/write shared state  
  
    // If testSharedState is now true  
    cv.signal(&lock);  
  
    // NO WARNING: signal keeps lock  
  
    // Read/write shared state  
    lock.release();  
}
```

# Condition Variables

- ALWAYS hold lock when calling wait, signal, broadcast
  - Condition variable is sync FOR shared state
  - ALWAYS hold lock when accessing shared state
- Condition variable is memoryless
  - If signal when no one is waiting, no op
  - If wait before signal, waiter wakes up
- Wait atomically releases lock
  - What if wait, then release?
  - What if release, then wait?

# Condition Variables, cont'd

- When a thread is woken up from wait, it may not run immediately
  - Signal/broadcast puts thread on ready list
  - When lock is released, anyone might acquire it
- Wait MUST be in a loop

```
while ( needToWait() ) {  
    condition.Wait(lock);  
}
```
- Simplifies implementation
  - Of condition variables and locks
  - Of code that uses condition variables and locks

# Structured Synchronization

- Identify objects or data structures that can be accessed by multiple threads concurrently
- Add locks to object/module
  - Grab lock on start to every method/procedure
  - Release lock on finish
- If need to wait
  - `while( needToWait() ) { condition.Wait(&lock); }`
  - Do not assume when you wake up that signaller just ran
- If do something that might wake someone up
  - Signal or Broadcast
- Always leave shared state variables in a consistent state
  - When lock is released, or when waiting

# Remember the rules

- Use consistent structure
- Always use locks and condition variables
- Always acquire lock at beginning of procedure, release at end
- Always hold lock when using a condition variable
- Always wait in while loop
- Never spin in sleep()

(if time permits)

# Pthread Syntax

## Thread B

## Thread A

```
pthread_mutex_lock( &my_lock );
```

```
while( /* test of shared state fails  
*/ ){
```

```
pthread_cond_wait( &my_cond_var,  
                  &my_lock );  
}
```

```
/* read/write shared state */
```

```
pthread_mutex_unlock( &my_lock );
```

```
pthread_mutex_lock( &my_lock );
```

```
/* read/write shared state */
```

```
/* if state has changed in a way  
that
```

```
    allows other threads to make  
    progress, then signal or  
broadcast */
```

```
pthread_cond_signal( &my_cond_  
var );
```

```
pthread_mutex_unlock( &my_loc  
k );
```



# Blocking Bounded Buffer in Java

```
class BoundedBuffer {  
    Lock lock = new ReentrantLock();  
    Condition notFull =  
lock.newCondition();  
    Condition notEmpty =  
lock.newCondition();  
    Object[] items = new Object[100];  
    int putptr, takeptr, count;  
  
    public void put(Object x) throws IE {  
        lock.lock(); try {  
            while ( count == items.length )  
                notFull.await();  
            items[putptr] = x;  
            if ( ++putptr == items.length )  
                putptr = 0;  
            ++count;  
            notEmpty.signal();  
        } finally { lock.unlock(); }  
    }  
}
```

```
    public Object take() throws IE {  
        lock.lock(); try {  
            while ( count == 0 )  
                notEmpty.await();  
            Object x = items[takeptr];  
            if ( ++takeptr ==  
items.length )  
                takeptr = 0;  
            --count;  
            notFull.signal();  
            return x;  
        } finally { lock.unlock(); }  
    }  
}
```

(code example from  
[gee.cs.oswego.edu/dl/concurrency-interest/jsr166-slides.pdf](http://gee.cs.oswego.edu/dl/concurrency-interest/jsr166-slides.pdf))

# Pthreads and Java 5 (and later)

When waiting upon a Condition Variable, a “spurious wakeup” is permitted to occur, in general, as a concession to the underlying platform semantics. This has little practical impact on most application programs as a Condition Variable should always be waited upon in a loop, testing the state predicate that is being waited for.

# Mesa vs. Hoare semantics

- Mesa
  - Signal puts waiter on ready list
  - Signaller keeps lock and processor
- Hoare
  - Signal gives processor and lock to waiter
  - When waiter finishes, processor/lock given back to signaller
  - Nested signals possible!

# FIFO Bounded Buffer (Hoare semantics)

```
get() {  
    lock.acquire();  
    if ( front == tail ) {  
        empty.wait(&lock);  
    }  
    item = buf[front % MAX];  
    front++; // ignoring ovf  
    full.signal(&lock);  
    lock.release();  
    return item;  
}
```

```
put(item) {  
    lock.acquire();  
    if ( (tail - front) == MAX ) {  
        full.wait(&lock);  
    }  
    buf[last % MAX] = item;  
    last++; // ignoring ovf  
    empty.signal(&lock);  
    // CAREFUL: someone else ran  
    lock.release();  
}
```

Initially: front = tail = 0; MAX is buffer capacity  
empty/full are condition variables

# FIFO Bounded Buffer (Mesa semantics)

- Create a condition variable for every waiter
- Queue condition variables (in FIFO order)
- Signal picks the front of the queue to wake up
- CAREFUL if spurious wakeups!
- Easily extends to case where queue is LIFO, priority, priority donation, ...

# FIFO Bounded Buffer

(Mesa semantics, put() is similar)

```
get() {  
    lock.acquire();  
    myPosition = numGets+  
    +;  
    // ignoring ovf  
    self = new Condition;  
    nextGet.append(self);  
    while ( front < myPosition  
    || front == tail ) {  
        self.wait(&lock);  
    }  
    delete self;  
    item = buf[front % MAX];  
    front++; // ignoring ovf  
    if (next =  
    nextPut.remove()) {  
        next->signal(&lock);  
    }  
    lock.release();  
    return item;  
}
```

Initially: front = tail = numGets = 0; MAX is buffer capacity  
nextGet, nextPut are queues of Condition Variables