

# Introduction to Operating Systems

CPSC/ECE 3220 Summer 2018

Lecture Notes  
OSPP Chapter 5 – Part B

(adapted by Mark Smotherman from Tom Anderson's slides on OSPP web site)

# Roadmap

Concurrent Applications

---

Shared Objects

Bounded Buffer

Barrier

---

Synchronization Variables

Semaphores

Locks

Condition Variables

---

Atomic Instructions

Interrupt Disable

Test-and-Set

---

Hardware

Multiple Processors

Hardware Interrupts

---

# Implementing Synchronization

Take 1: using memory load/store

- See too much milk solution/Peterson's algorithm

Take 2:

```
Lock::acquire()  
    { disable interrupts }  
Lock::release()  
    { enable interrupts }
```

# Lock Implementation, Uniprocessor

```
Lock::acquire() {  
    disableInterrupts();  
    if (value == BUSY) {  
        waiting.add(myTCB);  
        myTCB->state = WAITING;  
        next = readyList.remove();  
        switch(myTCB, next);  
        myTCB->state = RUNNING;  
    } else {  
        value = BUSY;  
    }  
    enableInterrupts();  
}
```

```
Lock::release() {  
    disableInterrupts();  
    if (!waiting.Empty()) {  
        next = waiting.remove();  
        next->state = READY;  
        readyList.add(next);  
    } else {  
        value = FREE;  
    }  
    enableInterrupts();  
}
```

# Multiprocessor

- Read-modify-write instructions
  - Atomically read a value from memory, operate on it, and then write it back to memory
  - Intervening instructions prevented in hardware
- Examples
  - Test and Set
  - Exchange (Intel: xchgb, w/ lock prefix to make atomic)
  - Compare and Swap
- Any of these can be used for implementing locks and condition variables!

# Spinlocks

A spinlock is a lock where the processor waits in a loop for the lock to become free

- Assumes lock will be held for a short time
- Used to protect the CPU scheduler and to implement locks

```
Spinlock::acquire() {  
    while ( TestAndSet(&lockValue) == BUSY )  
        ;  
}  
Spinlock::release() {  
    lockValue = FREE;  
    memorybarrier();  
}
```

# How many spinlocks?

- Various data structures
  - Queue of waiting threads on lock X
  - Queue of waiting threads on lock Y
  - List of threads ready to run
- One spinlock per kernel?
  - Bottleneck!
- Instead:
  - One spinlock per lock
  - One spinlock for the scheduler ready list
    - Perhaps per-core ready lists: one spinlock per core

# Lock Implementation, Multiprocessor

```
Lock::acquire() {  
    disableInterrupts();  
    spinLock.acquire();  
    if (value == BUSY) {  
        waiting.add(myTCB);  
        suspend(&spinlock);  
    } else {  
        value = BUSY;  
    }  
    spinLock.release();  
    enableInterrupts();  
}
```

```
Lock::release() {  
    disableInterrupts();  
    spinLock.acquire();  
    if (!waiting.Empty()) {  
        next = waiting.remove();  
        scheduler->  
makeReady(next);  
    } else {  
        value = FREE;  
    }  
    spinLock.release();  
    enableInterrupts();  
}
```



# Lock Implementation, Multiprocessor

```
Sched::suspend(SpinLock *lock) {  
    TCB *next;  
  
    disableInterrupts();  
    schedSpinLock.acquire();  
    lock->release();  
    myTCB->state = WAITING;  
    next = readyList.remove();  
    thread_switch(myTCB, next);  
    myTCB->state = RUNNING;  
    schedSpinLock.release();  
    enableInterrupts();  
}  
  
Sched::makeReady(TCB  
    *thread) {  
  
    disableInterrupts ();  
    schedSpinLock.acquire();  
    readyList.add(thread);  
    thread->state = READY;  
    schedSpinLock.release();  
    enableInterrupts();  
}
```

# What thread is currently running?

- Thread scheduler needs to find the TCB of the currently running thread
  - To suspend and switch to a new thread
  - To check if the current thread holds a lock before acquiring or releasing it
- On a uniprocessor, easy: just use a global
- On a multiprocessor, various methods:
  - Compiler dedicates a register (e.g., r31 points to TCB running on the this CPU; each CPU has its own r31)
  - If hardware has a special per-processor register, use it
  - Fixed-size stacks: put a pointer to the TCB at the bottom of its stack
    - Find it by masking the current stack pointer

# Lock Implementation, Linux

- Most locks are free most of the time
  - Why?
  - Linux implementation takes advantage of this fact
- Fast path
  - If lock is FREE, and no one is waiting, two instructions to acquire the lock
  - If no one is waiting, two instructions to release the lock
- Slow path
  - If lock is BUSY or someone is waiting, use multiprocessor impl.
- User-level locks
  - Fast path: acquire lock using test&set
  - Slow path: system call to kernel, use kernel lock

# Lock Implementation, Linux

```
struct mutex {  
    /* 1: unlocked ; 0:  
       locked; negative :  
       locked, possible  
       waiters */  
    atomic_t count;  
    spinlock_t wait_lock;  
    struct list_head  
        wait_list;  
};
```

```
// atomic decrement  
// %eax is pointer to  
    count  
lock decl (%eax)  
jns 1f // jump if not  
    signed  
        // (if value is  
        now 0)  
call slowpath_acquire  
1:
```

# Semaphores

- Semaphore has a non-negative integer value
  - P() atomically waits for value to become  $> 0$ , then decrements
  - V() atomically increments value (waking up waiter if needed)
- Semaphores are like integers except:
  - Only operations are P and V
  - Operations are atomic
    - If value is 1, two P's will result in value 0 and one waiter
- Semaphores are useful for
  - Unlocked wait: interrupt handler, fork/join

# Bounded Buffer with Locks/CVs

```
get() {  
    lock.acquire();  
    while (front == tail) {  
        empty.wait(&lock);  
    }  
    item = buf[front % MAX];  
    front++;  
    full.signal(&lock);  
    lock.release();  
    return item;  
}
```

```
put(item) {  
    lock.acquire();  
    while ((tail - front) ==  
MAX) {  
        full.wait(&lock);  
    }  
    buf[tail % MAX] = item;  
    tail++;  
    empty.signal(&lock);  
    lock.release();  
}
```

Initially: front = tail = 0; MAX is buffer capacity  
CVs: empty and full

# Bounded Buffer with Semaphores

```
get() {  
    fullSlots.P();  
    mutex.P();  
    item = buf[front %  
MAX];  
    front++;  
    mutex.V();  
    emptySlots.V();  
    return item;  
}
```

```
put(item) {  
    emptySlots.P();  
    mutex.P();  
    buf[last % MAX] =  
item;  
    last++;  
    mutex.V();  
    fullSlots.V();  
}
```

Initially: front = last = 0; MAX is buffer capacity  
Semaphores: mutex = 1; emptySlots = MAX; fullSlots

# Compare Semaphore P()/V() Implementation with Locks/CVs

```
Semaphore::P() {  
    disableInterrupts();  
    spinLock.acquire();  
    if (value == 0) {  
        waiting.add(myTCB);  
        suspend(&spinlock);  
    } else {  
        value--;  
    }  
    spinLock.release();  
    enableInterrupts();  
}
```

```
Semaphore::V() {  
    disableInterrupts();  
    spinLock.acquire();  
    if (!waiting.Empty()) {  
        next = waiting.remove();  
        scheduler->  
makeReady(next);  
    } else {  
        value++;  
    }  
    spinLock.release();  
    enableInterrupts();  
}
```



# Remember the rules

- Use consistent structure
- Always use locks and condition variables
- Always acquire lock at beginning of procedure, release at end
- Always hold lock when using a condition variable
- Always wait in while loop
- Never spin in sleep()

(if time permits)

# Communicating Sequential Processes (CSP/Google Go)

- A thread per shared object
  - Only thread allowed to touch object's data
  - To call a method on the object, send thread a message with method name, arguments
  - Thread waits in a loop, get msg, do operation
- No memory races!

# Locks/CVs vs. CSP

- Create a lock on shared data
  - = create a single thread to operate on data
- Call a method on a shared object
  - = send a message/wait for reply
- Wait for a condition
  - = queue an operation that can't be completed just yet
- Signal a condition
  - = perform a queued operation, now enabled

[ bounded\_buffer || producer || consumer ]

|| is  
concurrent  
execution

producer :: \*[ <produce item> ;  
bounded\_buffer ! item  
]

\*[ ] is  
repetition

consumer :: \*[ bounded\_buffer ? item;  
<consume item>  
]

! is  
send

bounded\_buffer :: buffer: (0..9) item;  
count, in, out: integer;

? is  
receive

count := 0; in := 0; out := 0;

\*[ count < 10 & producer ? buffer(in) ->  
in := (in + 1) mod 10;  
count := count + 1  
||

count > 0 & consumer ! buffer(out) ->  
out := (out + 1) mod 10;  
count := count - 1  
]

-> marks a  
guarded  
statement

# Implementing Condition Variables using Semaphores (Take 1)

```
wait(lock) {  
    lock.release();  
    semaphore.P();  
    lock.acquire();  
}  
signal() {  
    semaphore.V();  
}
```

# Implementing Condition Variables using Semaphores (Take 2)

```
wait(lock) {  
    lock.release();  
    semaphore.P();  
    lock.acquire();  
}  
signal() {  
    if (semaphore is not empty)  
        semaphore.V();  
}
```

# Implementing Condition Variables using Semaphores (Take 3)

```
wait(lock) {  
    semaphore = new Semaphore;  
    queue.Append(semaphore); // queue of waiting threads  
    lock.release();  
    semaphore.P();  
    lock.acquire();  
}  
signal() {  
    if (!queue.Empty()) {  
        semaphore = queue.Remove();  
        semaphore.V(); // wake up waiter  
    }  
}
```